

Ved Thakar

780-964-8081 | ved.thakar@mail.utoronto.ca | [linkedin.com/in/ved-thakar-00202b247](https://www.linkedin.com/in/ved-thakar-00202b247) | <https://vedthakar.github.io>

TECHNICAL SKILLS

GenAI & LLM Applications: LLM-Based Systems, Retrieval-Augmented Generation, Vector Search & Embeddings (pgvector), Prompt Engineering, Model Output Evaluation & Quality Review, Frontier Model Selection, Designing Around Model Failure, Gemini & Groq APIs

AI Agents & Experimentation: Agentic Orchestration, Multi-Agent Frameworks, MCP Integrations, Human-in-the-Loop Review Gates, Evaluation Harnesses & Benchmarking, Experimentation Infrastructure, Parallel Training Jobs (AWS SageMaker), Annotated & Curated Datasets

Product Analytics & Data: Product Metric Definition & Monitoring, Usage & Behavioural Data Analysis, Scoring & Ranking Models, Dashboards & Self-Service Reporting, SQL (PostgreSQL), Cloud Data Warehouses (BigQuery), ETL Pipelines, Time-Series Trend Analysis, pandas, NumPy, scikit-learn

Languages & Prototyping: Python, SQL, TypeScript, JavaScript, Java, GoLang, C++, C; React & Next.js for rapid prototyping, Tailwind CSS, FastAPI, Node.js, Visual Studio Code & Cursor, Git/GitHub

Documentation & Communication: Stakeholder-Facing Documentation (Notion, Markdown, Confluence-equivalent), Published Technical Writing, Architecture Notes & Decision Records, Presenting Findings to Technical & Non-Technical Audiences

Ways of Working: Entrepreneurial, Works with Minimal Guidance, Translates Ambiguous Tasks into Clear Outputs, Solves Own Problems, Curious About Where Product and ML Intersect

EXPERIENCE

Growth & Solutions Fellow

2026 – Present

Tavus (AI product company; YC, Sequoia & CRV-backed)

Remote / San Francisco

- Joined an AI company where no account list is handed to you. On my first trip to San Francisco I sat with a prospective customer and mapped what they were trying to accomplish onto how the AI product would actually be applied, helping close a **\$750K deal**.
- Sourced and developed opportunities myself rather than waiting on inbound, **adding millions in new pipeline**, and stayed close enough to those customers to carry what they needed back into product conversations. This is the product and ML intersection I keep gravitating toward.

Backend Software Engineer Intern

01/2026 to 04/2026

Geotab (Data Pipelines, Python, Analytics)

Toronto, ON

- Handed an ambiguous delivery-visibility problem with no defined scope, I decomposed it into clear deliverables before writing code, then built and operated the pipeline that fed the reporting, **improving delivery visibility by 80%** for the teams analyzing that data.
- The team had no fast feedback on regressions, so I built the automated test coverage and CI gating, **cutting cycle time by 70%**, and documented the architecture so stakeholders outside engineering could follow how it worked.

Junior Solutions Architect (Intern)

05/2025 to 03/2026

Savi Finance (GenAI Tooling, AWS, Automation)

Toronto, ON

- Found a manual process consuming developer hours that nobody had scoped or assigned, decided it was mine to solve without waiting for direction, and shipped an event-driven service on AWS Lambda and SQS, **cutting turnaround time by 80%**.
- Used GenAI coding tools daily on a production system but reviewed every generated change before merge, which is where I learned to treat model output as a draft whose quality I owned rather than a result to accept.

Software Developer Intern

May 2025 – Sept. 2025

Ascenix (early-stage startup, YC F26; React, TypeScript)

Remote

- At a Y Combinator-backed startup with nobody checking work hourly, carried customer-facing React and TypeScript features from ambiguous ticket to released product inside a single term.
- Rather than stopping once a feature worked, tested each against the API contract before handoff and fixed what review surfaced, so shipped work did not return to the team.

PROJECTS/COMPETITIONS

DealFlow OS — LLM Product with Analytics | *Python, SQL, React, PostgreSQL, LLMs*

- Investment firms were deciding from data scattered across sources, so I built and still operate a platform solo, now in production at **three professional services firms**, defining the usage and scoring metrics the product reports on.
- My first version let an LLM write the research memos and it produced fluent output nobody could verify, so I rebuilt it: memos assemble deterministically from source records with the model only refining wording, which means a rate limit degrades phrasing instead of inventing a figure.
- Analysed behavioural usage data to decide which signals actually indicated momentum and which only looked meaningful, then shipped dashboards on top across **500+ production deployments**.

SWAMR — AI Agent Framework | *TypeScript, Node.js, LLMs, MCP, Cursor*

- AI coding tools run one task at a time, so I built and published a framework coordinating **150+ specialist agents** in parallel across planning, build, test, and hardening phases, with MCP servers wired to GitHub, Supabase and Slack.
- Agent output cannot be trusted blindly, so I designed mandatory human review gates keeping generated work a draft until a person validates it, plus checkpointing that resumes an interrupted run exactly where it stopped.

Revenant — LLM Retrieval Platform, Winner, GenAI Genesis 2026 | *Python, FastAPI, pgvector*

- Teams lose knowledge when the person holding it leaves, so we co-built a system embedding GitHub, Slack and Jira activity into a vector index and answering natural language questions over it, **winning the sponsor prize for best use of memory architecture**.
- Batch processing was blocking live requests, so I designed the schema and query paths to separate a low-latency retrieval path from a batch evaluation path.

Model Evaluation & Benchmarking | *AWS SageMaker, Python, scikit-learn, Embeddings*

- Model comparisons at a **36-hour** research hackathon were being made on whichever run looked best, so I built an evaluation harness scoring every configuration on identical metrics, and ran parallel SageMaker training jobs to fit the experiment loop inside the deadline.
- Separately, a founder needed to pick a model without the background to compare options, so I evaluated frontier candidates and wrote the recommendation selecting Llama 3.1 8B on licensing and architectural grounds, in language she could act on.

TripBuddy — Acquired Product | *Product Ownership, Swift, Flutter, Django*

- I believed our problem was performance, so I took the product on a real trip and watched four first-time users fail at flows I was sure were intuitive. That proved the problem was comprehension, and I used the finding to justify a full rebuild rather than incremental fixes.
- Led four people from a **24-hour** build to a product **a startup acquired** two weeks after a live demo, with full codebase handover in June 2026.

LEADERSHIP & INVOLVEMENT

Hackathon Judge & Mentor	2025 – 2026
<i>Student Hackathons (2 events)</i>	<i>Toronto, ON</i>

- Judged submissions against a consistent rubric under a tight window, and mentored first-time builders, where the job is explaining something hard in plain language to someone who does not share your context.

EDUCATION

University of Toronto	Toronto, ON
<i>BSc Computer Science & Economics (Expected Graduation: Aug. 2028)</i>	<i>Sept. 2024 – Aug. 2028</i>

- Relevant coursework: Machine Learning (CSC311), Numerical Methods (CSC336), Data Structures in C (CSCA48), Microeconomics, Formal Logic, Linear Algebra.

Google Cloud	Toronto, ON
<i>Terraform for Google Cloud — Infrastructure as Code Certification</i>	<i>Feb. 2026 – Mar. 2026</i>